

PRJ 6

PRJ 6

一. Debug记录

TASK 1

- (1) 未pad镜像
- (2) 分配inode时犯的低级错误.....
- (3) 忽视目录项被回收后目录块中的空泡
- (4) 上板后的一些迷死

TASK 2

- (1) 读写文件的实际长度?
- (2) 多级索引考虑不周

TASK 3 [Undone]

3. 未完待续

二. Design Review

1. Q &A

2. Design Review后的改进

- 1) 元数据的个数&一致性
- 2) 删除时出现的问题
- 3) 一个南死人的bug: 内核栈溢出

一. Debug记录

• TASK 1

- (1) 未pad镜像

```

----- COMMAND -----
blocks read error!fs
[FS] Start initialize filesystem!
[FS] Setting superbblock...
  magic: 0x20221205
  num sector: 1049346, start sector: 1048576
  sector map offset: 1(256)
  inode map offset: 257(1)
blocks write error!1048576)
blocks write error!map...
blocks write error!-map...
blocks read error!...
blocks write error!
blocks read error!
blocks write error!
blocks write error!
blocks read error!
blocks write error!

```

开局不利.....意识到此时应该暂时给镜像pad一下，按照prj4修改Makefile:

```

1 image: $(ELF_CREATEIMAGE) $(ELF_BOOT) $(ELF_MAIN) $(ELF_USER)
2 #   cd $(DIR_BUILD) && ./$(<F) --extended $(filter-out $(<F), $(^F))
3     cd $(DIR_BUILD) && ./$(<F) --extended $(filter-out $(<F), $(^F))
  && dd if=/dev/zero of=image oflag=append conv=notrunc bs=512MB
  count=2

```

```

----- COMMAND -----
> root@UCAS_OS: mkfs
[FS] Start initialize filesystem!
[FS] Setting superbblock...
  magic: 0x20221205
  num sector: 1049346, start sector: 1048576
  block map offset: 1(256)
  inode map offset: 257(1)
  inode offset: 258(512)
  data offset: 770(1048576)
[FS] Setting inode-map...
[FS] Setting sector-map...
blocks write error!..
[FS] Initialize filesystem finished!
> root@UCAS_OS:

```

- (2) 分配inode时犯的低级错误.....

```

----- COMMAND -----
> root@UCAS_OS:~/ $ mkfs
[FS] Start initialize filesystem!
[FS] Setting superbblock...
    magic: 0x20221205
    num sector: 1049346, start sector: 1048576
    block map offset: 1(256)
    inode map offset: 257(1)
    inode offset: 258(512)
    data offset: 770(1048576)
[FS] Setting inode-map...
[FS] Setting sector-map...
[FS] Setting inode...
[FS] Initialize filesystem finished!
> root@UCAS_OS:~/ $ mkdir hey
> root@UCAS_OS:~/ $ ls
    hey
> root@UCAS_OS:~/ $ mkdir hi
> root@UCAS_OS:~/ $ ls
    hi
> root@UCAS_OS:~/ $

```

mkdir时似乎覆盖了之前创建的目录，gdb跟踪：

```

1 Breakpoint 1, do_mkdir (path=0x1269c "hey")
2   at ./kernel/fs/fs.c:238
3 238         if(get_inode_from_name(current_inode, path, NULL))
4 (gdb) n
5 242         int ino = alloc_inode();
6 (gdb)
7 244         bzero(buffer, 512);
8 (gdb)
9 245         dentry_t * de = (dentry_t*) buffer;
10 (gdb)
11 246         strcpy(de[0].name, ".");
12 (gdb)
13 247         strcpy(de[1].name, "..");
14 (gdb)
15 248         de[0].ino = ino;
16 (gdb)
17 249         de[1].ino = current_inode.ino;
18 (gdb)
19 250         uint32_t data_blk_addr = alloc_block();
20 (gdb)
21 251         bios_sdwrite(kva2pa(buffer), 1, data_blk_addr);
22 (gdb) p data_blk_addr
23 $1 = 131072
24 (gdb) n
25 253         inode_t *node = ino2inode(ino);
26 (gdb) p ino
27 $2 = 513
28 (gdb)
29 Breakpoint 1, do_mkdir (path=0x1269c "hi")
30   at ./kernel/fs/fs.c:238
31 238         if(get_inode_from_name(current_inode, path, NULL))
32 (gdb) n

```

```

33 242         int ino = alloc_inode();
34 (gdb)
35 244         bzero(buffer, 512);
36 (gdb) p ino
37 $3 = 513
38 (gdb)

```

意识到分配inode的函数出了问题：

```

1  for(i=0; i<INODE_MAP_NUM*SECTOR_SIZE; i++){
2      for(j=0, mask=1; j<sizeof(char); j++, mask<<1){
3          if(imap[i] & mask == 0)
4              break;
5      }
6  }
7  // 将未使用的inode结点置为已用
8  imap[i] |= mask;
9  bios_sdwrite(kva2pa(imap), INODE_MAP_NUM, FS_START_SEC +
    INODE_MAP_OFFSET);

```

意识到break只跳出了一重循环，同时注意到按位与的优先级低于==。

- (3) 忽视目录项被回收后目录块中的空泡

修改后可以正常创建目录，但删除文件夹会出错：

```

----- COMMAND -----
> root@UCAS_OS:~/$ mkfs
[FS] Start initialize filesystem!
[FS] Setting superblock...
    magic: 0x20221205
    num sector: 1049346, start sector: 1048576
    block map offset: 1(256)
    inode map offset: 257(1)
    inode offset: 258(512)
    data offset: 770(1048576)
[FS] Setting inode-map...
[FS] Setting sector-map...
[FS] Setting inode...
[FS] Initialize filesystem finished!
> root@UCAS_OS:~/$ mkdir hi
> root@UCAS_OS:~/$ mkdir hey
> root@UCAS_OS:~/$ ls
    hi hey
> root@UCAS_OS:~/$ rmdir hi
> root@UCAS_OS:~/$ ls
> root@UCAS_OS:~/$

```

原因是起初遇到空结点就会跳出循环：

```

// 跳过.和..
for(int i=2; i<SECTOR_SIZE/sizeof(dentry_t); i++){
    if(de[i].name[0]==0)
        break;
}

```

而对于位于其后的结点就可能无法被遍历，修改为continue后就可正常进行。

- (4) 上板后的一些迷死

发现在ls-l时总是会卡住，查看实现：

```
1  for(int i=2; i<DPSEC; i++){
2      if(de[i].name[0]==0)
3          continue;
4      else if(option){ // 需打印详细信息
5          inode_t tmp = *ino2inode(de[i].ino);
6          printk("%c%c%c nlink:%d ctime:%d atime:%d mtime:%d
size:%d %s\n",
7              tmp.type == T_DIR ? 'd' : '-',
8              tmp.mode & O_RDONLY ? 'r' : '-',
9              tmp.mode & O_WRONLY ? 'w' : '-',
10             tmp.nlink, tmp.ctime, tmp.atime, tmp.mtime,
11             tmp.size,
12             de[i].name
13             );
14     }
15     else
16         printk("\t%s", de[i].name);
17 }
```

原本是在读到名字为空时就不会根据其ino域去获取inode，但由于sd卡中可能原本就有些脏数据，导致此时可能会获取到很诡异的ino。故需要在每次分配block时先将数据清空。

• TASK 2

- (1) 读写文件的实际长度？

运行测试用例卡住，gdb跟踪：

```
1  (gdb) n
2  608          if(write_ptr % BLOCK_SIZE)
3  <p write_ptr
4  $3 = 0
5  (gdb) n
6  610          memcpy(buffer + (write_ptr%BLOCK_SIZE),
7             buff,partial_len);
8  (gdb)
9  ^C
10 Program received signal SIGINT, Interrupt.
11 atomic_swap (
```

```

11     val=1,
12     mem_addr=18446743800176394880) at
    ./arch/riscv/include/atomic.h:9
13 9      {
14  (gdb)

```

意识到此时len的设置不准确：

```

1  // 以block为单位读取
2  for(int read_ptr = fdesc_array[fd].read_ptr;
   read_ptr < fdesc_array[fd].read_ptr + len;){
3      int partial_len = read_ptr % BLOCK_SIZE ? (BLOCK_SIZE -
   (read_ptr % BLOCK_SIZE)) : BLOCK_SIZE;
4      uint32_t read_addr = get_data_block_addr(node, read_ptr);
5      bios_sdread(kva2pa(buffer), BLOCK_SIZE/SECTOR_SIZE,
   read_addr);
6      memcpy(buff, buffer+(read_ptr%BLOCK_SIZE), partial_len);
7      read_ptr += partial_len;
8      buff += partial_len;
9  }

```

还需要判断 `partial_len` 和要求读入的 `len` 的关系：

```

1  int partial_len = read_ptr % BLOCK_SIZE ? (BLOCK_SIZE - (read_ptr %
   BLOCK_SIZE)) : BLOCK_SIZE;
2      int tmp = fdesc_array[fd].read_ptr + len -
   fdesc_array[fd].read_ptr;
3      partial_len = partial_len > tmp ? tmp : partial_len;

```

修改后可正常输出（请忽略间歇性抽风的终端输出）：

```

hello world!  c          v          : 00000004
hello world!
hello world!
hello world!
hello world!
hello world!
hello world!
hello world!
hello world!
hello world!
hello world!
-----[12[[1[11;25H-----
> root@UCAS_OS:~$ e ec rwfile &
Info: excute rwfile successfully, pid = 2
> root@UCAS_OS:~$

```

重启后调用cat也能正常输出：

```

.. -----[11H1;23H -----
> root@UCAS_OS:~$ cat 1.txt
hello world!
hello world!
hello world!
hello world!
hello world!
hello world!
hello world!
hello world!
hello world!
hello world!
QEMU: Terminated

```

- (2) 多级索引考虑不周

同时为了测试大文件的情况，在测试用例中增加lseek:

```

1 //-----*****REVISE START*****-----
2
3     assert(fd>=0);
4     assert(sys_lseek(fd, RW_START, SEEK_SET)==RW_START);
5 //-----*****REVISE END*****-----

```

```

e " * t2hellocworld!      v      : 00000004
hello world!
hello world!
hello world!
hello world!
hello world!
hello world!
hello world!
hello world!

```

```

----- COMMAND -----
> root@UCAS_OS:~$ mkfs
[FS] Start initialize filesystem!
[FS] Setting superblock...
    magic: 0x20221205
    num sector: 1049346, start sector: 1048576
    block map offset: 1(256)
    inode map offset: 257(1)
    inode offset: 258(512)
    data offset: 770(1048576)
[FS] Setting inode-map...
[FS] Setting sector-map...
[FS] Setting inode...
[FS] Initialize filesystem finished!
> root@UCAS_OS:~$ touch 1.txt
> root@UCAS_OS:~$ exec rwfile &
Info: excute rwfile successfully, pid = 2
> root@UCAS_OS:~$

```

发现输出不对。gdb跟踪:

1 (gdb) p tmp

```

2  $2 = 13
3  (gdb) n
4  636          uint32_t write_addr = get_data_block_addr(node,
write_ptr);
5  (gdb) n
6  637          if(write_ptr % BLOCK_SIZE ||
partial_len < BLOCK_SIZE)
7  (gdb) p write_addr
8  $3 = 1057458
9  (gdb) p write_ptr
10 $4 = 8388608
11 (gdb) p/x write_ptr
12 $5 = 0x800000
13 (gdb) n
14 638          bios_sdread(kva2pa(buffer),
BLOCK_SIZE/SECTOR_SIZE, write_addr);
15 .....
16
17 Breakpoint 2, do_fread (fd=0,
18     buff=0x117f0 "", length=13)
19     at ./kernel/fs/fs.c:602
20 602          for(int read_ptr = fdesc_array[fd].read_ptr;
read_ptr < fdesc_array[fd].read_ptr + len;){
21 (gdb) n
22 603          int partial_len = read_ptr % BLOCK_SIZE ?
(BLOCK_SIZE - (read_ptr % BLOCK_SIZE)) : BLOCK_SIZE;
23 (gdb)
24 604          int tmp = fdesc_array[fd].read_ptr + len -
read_ptr;
25 (gdb)
26 605          partial_len = partial_len > tmp ? tmp :
partial_len;
27 (gdb)
28 606          uint32_t read_addr = get_data_block_addr(node,
read_ptr);
29 (gdb)
30 607          bios_sdread(kva2pa(buffer), BLOCK_SIZE/SECTOR_SIZE,
read_addr);
31 (gdb) p read_addr
32 $15 = 1049362
33 (gdb) p read_ptr
34 $16 = 8388608
35 (gdb) p/x read_ptr
36 $17 = 0x800000
37 (gdb)

```

发现前后分配的地址不对，查看原先多级索引的实现：

```
1 // 使用一级索引
```



```

2     else if(size < DIRECT_SIZE + INDIRECT_1ST_SIZE){
3         size -= DIRECT_SIZE;
4         int index1 = size/(BLOCK_SIZE*IA_PER_BLOCK);
5         int index2 = (size -
index1*BLOCK_SIZE*IA_PER_BLOCK)/BLOCK_SIZE;
6         if(node.indirect_addrs_1st[index1]==0){
7             uint32_t data_blk_addr = alloc_block();
8             setup_level_index(data_blk_addr, 0);
9             inode_t* node_ptr = ino2inode(node.ino);
10            node_ptr->indirect_addrs_1st[index1] = data_blk_addr;
11            int offset = node.ino / IPSEC;
12            bios_sdwrite(kva2pa(buffer), 1, FS_START_SEC +
INODE_OFFSET + offset);
13            bios_sdread(kva2pa(buffer), BLOCK_SIZE/SECTOR_SIZE,
node_ptr->indirect_addrs_1st[index1]);
14            uint32_t* addr_array = buffer;
15            return addr_array[index2];
16        }
17    }
18    // 使用二级索引
19    else if(size < DIRECT_SIZE + INDIRECT_1ST_SIZE +
INDIRECT_2ND_SIZE){
20        size -= (DIRECT_SIZE + INDIRECT_1ST_SIZE);
21        int index1 = size/(BLOCK_SIZE*IA_PER_BLOCK*IA_PER_BLOCK);
22        int index2 = (size -
index1*BLOCK_SIZE*IA_PER_BLOCK*IA_PER_BLOCK)/(BLOCK_SIZE*IA_PER_BLOCK);
23        int index3 = (size -
index1*BLOCK_SIZE*IA_PER_BLOCK*IA_PER_BLOCK -
index2*IA_PER_BLOCK)/BLOCK_SIZE;
24        if(node.indirect_addrs_2nd[index1]==0){
25            uint32_t data_blk_addr = alloc_block();
26            setup_level_index(data_blk_addr, 1);
27            inode_t* node_ptr = ino2inode(node.ino);
28            node_ptr->indirect_addrs_2nd[index1] = data_blk_addr;
29            int offset = node.ino / IPSEC;
30            bios_sdwrite(kva2pa(buffer), 1, FS_START_SEC +
INODE_OFFSET + offset);
31            // 获取二级索引项
32            bios_sdread(kva2pa(buffer), BLOCK_SIZE/SECTOR_SIZE,
data_blk_addr);
33            uint32_t* addr_array = buffer;
34            data_blk_addr = addr_array[index2];
35            // 获取一级索引项
36            bios_sdread(kva2pa(buffer), BLOCK_SIZE/SECTOR_SIZE,
data_blk_addr);
37            return addr_array[index3];
38        }
39    }

```

```

40     // 使用三级索引
41     else if(size < DIRECT_SIZE + INDIRECT_1ST_SIZE +
INDIRECT_2ND_SIZE + INDIRECT_3RD_SIZE){
42         size -= (DIRECT_SIZE + INDIRECT_1ST_SIZE +
INDIRECT_2ND_SIZE);
43         int index1 = size/(BLOCK_SIZE*IA_PER_BLOCK*IA_PER_BLOCK);
44         int index2 = (size -
index1*BLOCK_SIZE*IA_PER_BLOCK*IA_PER_BLOCK)/(BLOCK_SIZE*IA_PER_BLOCK);
45         if(node.indirect_addrs_3rd==0){
46             uint32_t data_blk_addr = alloc_block();
47             setup_level_index(data_blk_addr, 1);
48             inode_t* node_ptr = ino2inode(node.ino);
49             node_ptr->indirect_addrs_3rd = data_blk_addr;
50             int offset = node.ino / IPSEC;
51             bios_sdwrite(kva2pa(buffer), 1, FS_START_SEC +
INODE_OFFSET + offset);
52             // 获取二级索引项
53             bios_sdread(kva2pa(buffer), BLOCK_SIZE/SECTOR_SIZE,
data_blk_addr);
54             uint32_t* addr_array = buffer;
55             data_blk_addr = addr_array[index1];
56             // 获取一级索引项
57             bios_sdread(kva2pa(buffer), BLOCK_SIZE/SECTOR_SIZE,
data_blk_addr);
58             return addr_array[index2];
59         }
60     }

```

意识到上述都忘记考虑索引项已经分配的情况，应该修改为如下（只附上一级索引的代码，二级、三级同理）：

```

1     // 使用一级索引
2     else if(size < DIRECT_SIZE + INDIRECT_1ST_SIZE){
3         size -= DIRECT_SIZE;
4         int index1 = size/(BLOCK_SIZE*IA_PER_BLOCK);
5         int index2 = (size -
index1*BLOCK_SIZE*IA_PER_BLOCK)/BLOCK_SIZE;
6         uint32_t data_blk_addr;
7         if(node.indirect_addrs_1st[index1]==0){
8             data_blk_addr = alloc_block();
9             setup_level_index(data_blk_addr, 0);
10            inode_t* node_ptr = ino2inode(node.ino);
11            node_ptr->indirect_addrs_1st[index1] = data_blk_addr;
12            int offset = node.ino / IPSEC;
13            bios_sdwrite(kva2pa(buffer), 1, FS_START_SEC +
INODE_OFFSET + offset);
14        }

```

```

15         else
16             data_blk_addr = node.indirect_addrs_1st[index1];
17             bios_sdread(kva2pa(buffer), BLOCK_SIZE/SECTOR_SIZE,
data_blk_addr);
18             uint32_t* addr_array = buffer;
19             return addr_array[index2];
20     }

```

修改后如下：

```

hello world!  c          v          : 0000 error  Aa _ab_ * 1 of 1  ↑ ↓ ×
hello world!
hello world!
hello world!
hello world!
hello world!
hello world!
hello world!
hello world!
hello world!
hello world!
----- COMMAND -----
> root@UCAS_OS:~$ mkfs
[FS] Start initialize filesystem!
[FS] Setting superblock...
    magic: 0x20221205
    num sector: 1049346, start sector: 1048576
    block map offset: 1(256)
    inode map offset: 257(1)
    inode offset: 258(512)
    data offset: 770(1048576)
[FS] Setting inode-map...
[FS] Setting sector-map...
[FS] Setting inode...
[FS] Initialize filesystem finished!
> root@UCAS_OS:~$ touch 1.txt
> root@UCAS_OS:~$ exec rwfile &
Info: excute rwfile successfully, pid = 2
> root@UCAS_OS:~$

```

虽然在输出上一模一样，但是我们可以通过查看block的使用情况来判断索引的分配情况：

```

----- COMMAND -----
> root@UCAS_OS:~$ statfs
[FS] state:
    magic: 0x20221205
    total sector: 1049346, start sector: 1048576(00100000)
    block map offset: 1, occupied sector: 256
    inode map offset: 257, occupied sector: 1
    inode offset: 258, usage 2/512
    data offset: 770, usage 1027/1048576
    inode size: 100B, dir entry size: 20B
> root@UCAS_OS:~$

```

为建立多级索引，其会使用多个数据块。此外，还可以通过ls -l查看详细信息：

```

> root@UCAS_OS:~$ ls -l
-rw nlink:01 ctime:14 atime:14 mtime:14 size:8388738 1.txt
> root@UCAS_OS:~$

```

可见其size符合预期值。但注意到读写时还未更新时间戳信息，修改后效果如下：

```

----- COMMAND -----
> root@UCAS_OS:~$ ls -l
-rw nlink:1 ctime:28 atime:41 mtime:39 size:8388738 1.txt

```

简名	全名	中文	作用
atime	Access Time	访问时间	最后一次访问文件（读取或执行）的时间
ctime	Change Time	变化时间	最后一次改变文件（属性或权限）或者目录（属性或权限）的时间
mtime	Modify Time	修改时间	最后一次修改文件（内容）或者目录（内容）的时间

回收不完全：

```

----- COMMAND -----
> root@UCAS_OS:~$ rm 1.txt
> root@UCAS_OS:~$ statfs
[FS] state:
  magic: 0x20221205
  total sector: 1049346, start sector: 1048576(00100000)
  block map offset: 1, occupied sector: 256
  inode map offset: 257, occupied sector: 1
  inode offset: 258, usage 1/512
  data offset: 770, usage 899/1048576
  inode size: 100B, dir entry size: 20B
> root@UCAS_OS:~$

```

gdb查看回收后的bmap情况：

```

(gdb) x/1000x bmap
0xffffffffc050250520 <bmap>:      0xfefefeff      0xfefefefe      0xfefefefe      0xfefefefe
0xffffffffc050250530 <bmap+16>:  0xfefefefe      0xfefefefe      0xfefefefe      0xfefefefe
0xffffffffc050250540 <bmap+32>:  0xfefefefe      0xfefefefe      0xfefefefe      0xfefefefe
0xffffffffc050250550 <bmap+48>:  0xfefefefe      0xfefefefe      0xfefefefe      0xfefefefe
0xffffffffc050250560 <bmap+64>:  0xfefefefe      0xfefefefe      0xfefefefe      0xfefefefe
0xffffffffc050250570 <bmap+80>:  0xfefefefe      0xfefefefe      0xfefefefe      0xfefefefe
0xffffffffc050250580 <bmap+96>:  0xfefefefe      0xfefefefe      0xfefefefe      0xfefefefe
0xffffffffc050250590 <bmap+112>: 0xfefefefe      0xfefefefe      0xfefefefe      0xfefefefe
0xffffffffc0502505a0 <bmap+128>: 0x00000000      0x00000000      0x00000000      0x00000000
0xffffffffc0502505b0 <bmap+144>: 0x00000000      0x00000000      0x00000000      0x00000000
0xffffffffc0502505c0 <bmap+160>: 0x00000000      0x00000000      0x00000000      0x00000000
0xffffffffc0502505d0 <bmap+176>: 0x00000000      0x00000000      0x00000000      0x00000000

```

查看回收部分代码：

```

1 | int bno = data_blk_addr - FS_START_SEC - DATA_OFFSET;
2 | bmap[bno/8] &= ~(1 << (bno%8));

```

意识到回收时忽略了分配是以块为单位进行的，故bno应该修改为：

```

1 | int bno = (data_blk_addr - FS_START_SEC - DATA_OFFSET) *SECTOR_SIZE
  | / BLOCK_SIZE;
2 | bmap[bno/8] &= ~(1 << (bno%8));

```

修改后可以正常删除（请再次忽略抽风的终端）：

```

-----[11;22HH[11;26H-----
> root@UCAS_OS:~$ [1212;22H
> root@UCAS_OS:~$ statfs
[FS] state:
  magic: 0x20221205
  total sector: 1049346, start sector: 1048576(00100000)
  block map offset: 1, occupied sector: 256
  inode map offset: 257, occupied sector: 1
  inode offset: 258, usage 1/512
  data offset: 770, usage 1/1048576
  inode size: 100B, dir entry size: 20B
> root@UCAS_OS:~$

```

• TASK 3 [Undone]

小程序发送包时显示如下：

```

stu@stu:/$ cd ~
stu@stu:~$ cd p6-pktRxTx_elf/
stu@stu:~/p6-pktRxTx_elf$ sudo ./pktRxTx -m 1 -t 5
[sudo] password for stu:
Info: pktRxTx was built at Dec 17 2022, 07:21:43
1. tap0 (No Description Available)
2. br0 (No Description Available)
3. enp0s3 (No Description Available)
4. lo (No Description Available)
5. Pseudo-device that captures on all interfaces
6. Bluetooth Linux Monitor
7. Linux netfilter log (NFLOG) interface
8. Linux netfilter queue (NFQUEUE) interface
Enter the interface number (1-8): 1
Info: Here, MAC Address is 128:250:91:51:86:239, listening on device tap0 ...
Info: MAC Address of the opposite is 00:10:53:00:30:83 ...
Info: Input command. For example:
    --- 'send 60': send 60 packets to the opposite
    --- 'test 60': keep sending packets in 60 seconds
    --- 'quit': quit this program
> send
> Please input file name: plus.elf
> Opening 'plus.elf'.
> Size is 348
[INFO] sending head packet, size = 8B.
[INFO] sending 0 packet, size = 34B.
>
> quit
Info: Sender finishes its task and exits!
stu@stu:~/p6-pktRxTx_elf$ ls

```

我个亲娘欸这包都没发两个咋触发网卡中断啊？难不成多发几个废包？

• 3. 未完待续

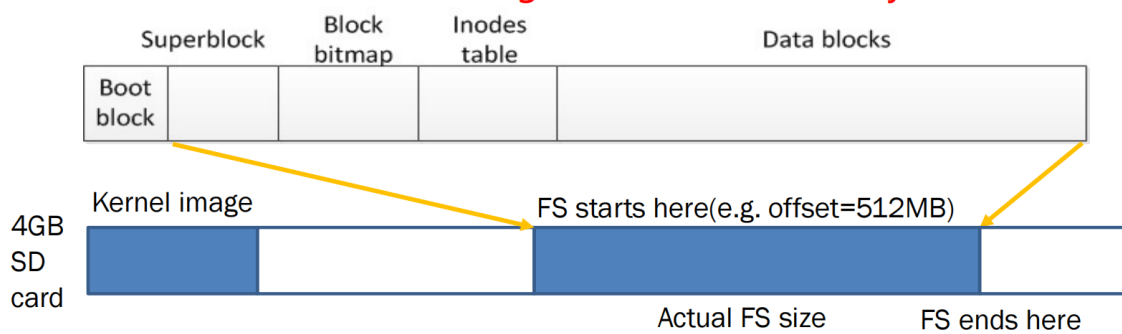
目前还需考虑如下问题：

- 考虑发包未满的情况？
- 考虑test文件在kernel内的处理方式：
 - 在task info数组里加上相关信息？
 - elf头是否要仿照createimage中处理？

二. Design Review

• 1. Q &A

- What is the disk layout in your design?



- Show the structures of your FS metadata, including superblock, inode, dentry, and file descriptor

```
1  /* data structures of file system */
2  typedef struct superblock_t{
3      // TODO [P6-task1]: Implement the data structure of
4      // superblock
5      uint32_t magic_number;
6      uint32_t fs_start_sec;
7      uint32_t fs_size;      // Size of file system image
8      (blocks)
9      uint32_t block_map_offset;
10     uint32_t block_map_num;
11     uint32_t inode_map_offset;
12     uint32_t inode_map_num;
13     uint32_t inode_offset;
14     uint32_t inode_num;
15     uint32_t data_offset;
16     uint32_t data_num;
17 } superblock_t;
```

```

17 typedef struct dentry_t{
18     // TODO [P6-task1]: Implement the data structure of
    directory entry
19     char name[16];
20     int ino;
21 } dentry_t;
22
23 typedef struct inode_t{
24     // TODO [P6-task1]: Implement the data structure of inode
25     char type;
26     char mode;
27     short nlink;          // Number of links to inode in file
    system
28     uint32_t ino;
29     uint32_t ctime;
30     uint32_t atime;
31     uint32_t mtime;
32     uint32_t size;
33     uint32_t direct_addrs[NDIRECT];
34     uint32_t indirect_addrs_1st[3];
35     uint32_t indirect_addrs_2nd[2];
36     uint32_t indirect_addrs_3rd;
37 } inode_t;
38
39 typedef struct fdesc_t{
40     // TODO [P6-task2]: Implement the data structure of file
    descriptor
41     uint8_t valid;
42     uint8_t mode;
43     short ref; // reference count
44     int ino;
45     uint32_t write_ptr;
46     uint32_t read_ptr;
47 } fdesc_t;
48
49

```

⋮ 着重讲一下nlink和ref用处区别

- How many files and directories do you file system support?

```

1 | #define IA_PER_BLOCK (BLOCK_SIZE/sizeof(uint32_t))
2 | #define DIRECT_SIZE (NDIRECT*BLOCK_SIZE)
3 | #define INDIRECT_1ST_SIZE (3*BLOCK_SIZE*IA_PER_BLOCK)
4 | #define INDIRECT_2ND_SIZE
   | (2*BLOCK_SIZE*IA_PER_BLOCK*IA_PER_BLOCK)
5 | #define INDIRECT_3RD_SIZE
   | (1*BLOCK_SIZE*IA_PER_BLOCK*IA_PER_BLOCK*IA_PER_BLOCK)
6 | #define MAX_FILE_SIZE (DIRECT_SIZE + INDIRECT_1ST_SIZE +
   | INDIRECT_2ND_SIZE + INDIRECT_3RD_SIZE)

```

一个inode最多：

- 索引 `MAX_FILE_SIZE` 大小的文件；
- 含 `MAX_FILE_SIZE/sizeof(dentry_t)` 个目录项；

文件系统使用 `256KB` 作inode的存储空间，故最多 `256K/sizeof(inode_t)` 个inode。

- What do you do when initializing a file system?
 - superblock的初始化
 - ⋮ 从磁盘特定位置载入后先check是否已建立fs
 - inode和block map的 初始化
 - ⋮ 数据块以4KB为粒度
 - 创建根目录
 - 分配inode
 - 分配目录索引页
 - 初始化文件描述符
- Given an operation, for example `ls /home/student`, How do you handle path lookup?
 - 路径解析：以 `/` 为分割
 - 相对路径特性：递归查找对应inode并返回

• 2. Design Review后的改进

- 1) 元数据的个数&一致性

在design review中意识到此前自己对元数据个数的理解尚不清晰，出现了inode数和inode map的体量不匹配的情况（原先是划定给inode的内存大小，反推inode数和inode map，后成功把自己算晕）。

后将inode数目划定，再利用相对关系定出inode map和inode内存大小，如下：

```
1 #define INODE_NUM          512          // INODE个数
2 #define DATA_BLOCK_NUM    (1<<20)     // 数据块个数（4KB为单位），共
    4GB
3 #define INODE_MAP_SEC_NUM  CEIL_DIV(INODE_NUM, SECTOR_SIZE*8)
    // inode map所占sector个数
4 #define BLOCK_MAP_SEC_NUM  CEIL_DIV(DATA_BLOCK_NUM, SECTOR_SIZE*8)
    // data block map所占sector个数
5 #define INODE_SEC_NUM      CEIL_DIV(sizeof(inode_t)*INODE_NUM,
    SECTOR_SIZE)          // inode占据的sector个数
6 #define DATA_BLOCK_SEC    (DATA_BLOCK_NUM/SECTOR_SIZE*BLOCK_SIZE)
    // data block占据的sector个数
```

- 2) 删除时出现的问题

- 硬链接文件的删除情况：原本设置为支持多层删除，后续发现删除了前一层文件后，本层目录下的硬链接文件反而会被删除，原因是根据name索引到了文件后我返回的是ino号，然后再在本层目录下以ino号做匹配删除。只需改为严格限制为单级删除即可。
- nlink数忘记更改：原本只是将之在内存中的缓存的nlink数减一，忘记写回。

- 3) 一个南死人的bug：内核栈溢出

在qemu时跑大文件都没有出问题，但是上板一直会卡住，后来擦边改了很多无关紧要的代码.....最后发现是函数开了过大的局部变量做buffer，在递归建立多级索引的时候把内核栈搞溢出了。后修改为每层递归使用不同的buffer即可：

```
1 static char level_buffer[3][BLOCK_SIZE];
2 void setup_level_index(uint32_t data_blk_addr, int level){
3     uint32_t* addr_array = level_buffer[level];
4     // printk("I am in setup_level_index\n");
5     alloc_block(addr_array, IA_PER_BLOCK);
6     if(level)
7         for(int i=0; i<IA_PER_BLOCK; i++)
8             setup_level_index(addr_array[i], level-1);
9     bios_sdwrite(kva2pa(level_buffer[level]),
    BLOCK_SIZE/SECTOR_SIZE, data_blk_addr);
10 }
```

- 很怪的是之前也开了一些4KB的局部buffer（我们的内核栈大小就4KB.....），但上板也没跑出问题。大概板子想出错时才出错，怪错只能自己抗🤡。

